

ChandraAILabs - Technical Architecture Guide

Platform Overview

Our enterprise HR RAG platform uses hybrid retrieval combining BM25 sparse search and Vertex AI Vector Sea

Components:

1. Gemini 2.5 Flash - Answer generation
2. Vertex AI Vector Search - Semantic retrieval
3. BM25 - Keyword retrieval
4. Cloud Run - Serverless deployment
5. Firestore - Chunk metadata store

Performance:

Policy queries: 3-5 seconds average

Personal queries: 2 seconds average

Cache hits: 0ms response time